



LOSS及其梯度

主讲人：龙良曲

Typical Loss

- Mean Squared Error
 - Cross Entropy Loss
 - binary
 - multi-class
 - +softmax
 - Leave it to Logistic Regression Part
-

MSE

- $\text{loss} = \sum [y - (xw + b)]^2$


- $L2 - \text{norm} = ||y - (xw + b)||_2$

- $\text{loss} = \text{norm}(y - (xw + b))^2$

Derivative

- $\text{loss} = \sum [y - f_{\theta}(x)]^2$

- $\frac{\nabla \text{loss}}{\nabla \theta} = 2 \sum [y - f_{\theta}(x)] * \frac{\nabla f_{\theta}(x)}{\nabla \theta}$



autograd.grad



```
In [15]: x=torch.ones(1)
In [17]: w=torch.full([1],2)
In [19]: mse=F.mse_loss(torch.ones(1), x*w)
Out[20]: tensor(1.)
```

```
In [21]: torch.autograd.grad(mse, [w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn
```

```
In [22]: w.requires_grad_()
Out[22]: tensor([2.], requires_grad=True)
```

```
In [23]: torch.autograd.grad(mse, [w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn
```

```
In [24]: mse=F.mse_loss(torch.ones(1), x*w)
```

```
In [25]: torch.autograd.grad(mse, [w])
Out[25]: (tensor([2.]),)
```

loss.backward



```
In [15]: x=torch.ones(1)
In [17]: w=torch.full([1],2)
In [19]: mse=F.mse_loss(torch.ones(1), x*w)
Out[20]: tensor(1.)
```

```
In [21]: torch.autograd.grad(mse,[w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn
```

```
In [22]: w.requires_grad_()
Out[22]: tensor([2.], requires_grad=True)
```

```
In [23]: torch.autograd.grad(mse,[w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn
```

```
In [24]: mse=F.mse_loss(torch.ones(1), x*w)
In [27]: mse.backward()
```

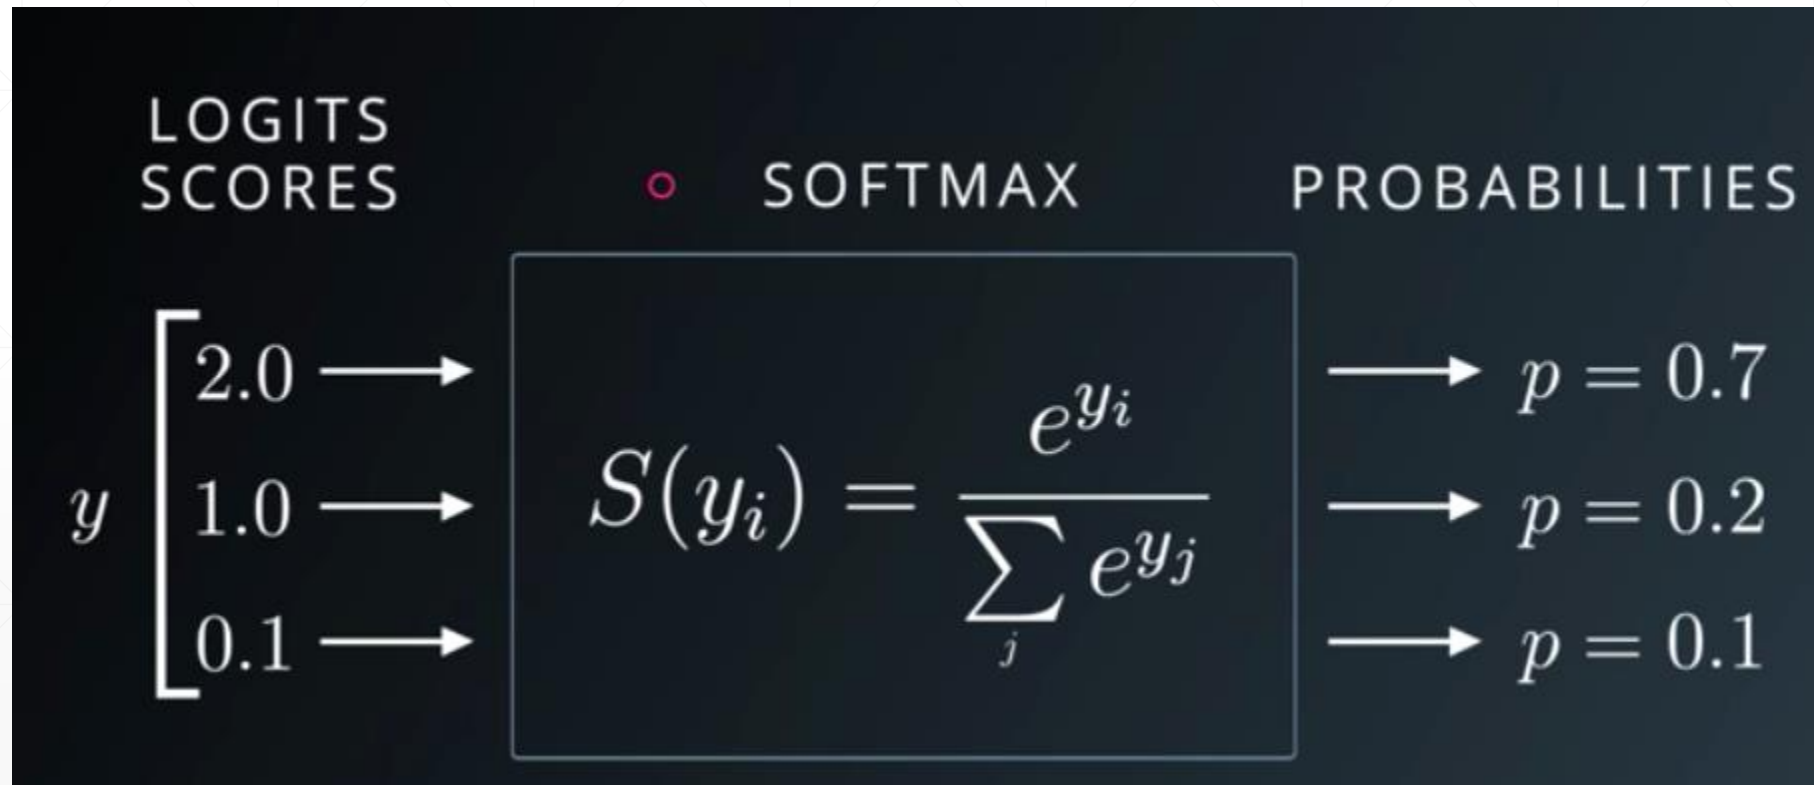
```
In [28]: w.grad
Out[28]: tensor([2.])
```

Gradient API

- `torch.autograd.grad(loss, [w1, w2,...])`
 - `[w1.grad, w2.grad...]`
 - `loss.backward()`
 - `w1.grad`
 - `w2.grad`
-

Softmax

- soft version of max



Derivative

导数

$$\frac{\partial p_i}{\partial a_j} = \frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j}$$

$$f(x) = \frac{g(x)}{h(x)}$$

$$f'(x) = \frac{g'(x)h(x) - h'(x)g(x)}{h(x)^2}$$

$$g(x) = e^{a_i}$$

$$h(x) = \sum_{k=1}^N e^{a_k}$$

$$p_i = \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}$$

when $i = j$

$$\begin{aligned} \frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j} &= \frac{e^{a_i} \sum_{k=1}^N e^{a_k} - e^{a_j} e^{a_i}}{\left(\sum_{k=1}^N e^{a_k}\right)^2} \\ &= \frac{e^{a_i} \left(\sum_{k=1}^N e^{a_k} - e^{a_j}\right)}{\left(\sum_{k=1}^N e^{a_k}\right)^2} \\ &= \frac{e^{a_j}}{\sum_{k=1}^N e^{a_k}} \times \frac{\left(\sum_{k=1}^N e^{a_k} - e^{a_j}\right)}{\sum_{k=1}^N e^{a_k}} \\ &= p_i (1 - p_j) \end{aligned}$$

Derivative

$$p_i = \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}$$

$$\frac{\partial p_i}{\partial a_j} = \frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j}$$

$$f(x) = \frac{g(x)}{h(x)}$$

$$f'(x) = \frac{g'(x)h(x) - h'(x)g(x)}{h(x)^2}$$

$$g(x) = e^{a_i}$$

$$h(x) = \sum_{k=1}^N e^{a_k}$$

when $i \neq j$

$$\frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j} = \frac{0 - e^{a_j} e^{a_i}}{\left(\sum_{k=1}^N e^{a_k}\right)^2}$$

$$\begin{aligned} &= \frac{-e^{a_j}}{\sum_{k=1}^N e^{a_k}} \times \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}} \\ &= -p_j \cdot p_i \end{aligned}$$

Derivative

$$\frac{\partial p_i}{\partial a_j} = \begin{cases} p_i(1 - p_j) & \text{if } i = j \\ -p_j \cdot p_i & \text{if } i \neq j \end{cases}$$

Or using Kronecker delta $\delta_{ij} = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$

$$\frac{\partial p_i}{\partial a_j} = p_i(\delta_{ij} - p_j)$$

F.softmax

```
In [29]: a=torch.rand(3) # tensor([0.1440, 0.5349, 0.7022])
```

```
In [33]: a.requires_grad_()
```

```
Out[33]: tensor([0.1440, 0.5349, 0.7022], requires_grad=True)
```

```
In [34]: p=F.softmax(a,dim=0)
```

```
In [35]: p.backward()
```

RuntimeError: Trying to backward through the graph a second time, but the buffers have already been freed. Specify retain_graph=True when calling backward the first time.

```
In [38]: p=F.softmax(a,dim=0)
```

```
In [39]: torch.autograd.grad(p[1],[a],retain_graph=True)
```

```
Out[39]: (tensor([-0.0828, 0.2274, -0.1447]),)
```

```
In [40]: torch.autograd.grad(p[2],[a])
```

```
Out[40]: (tensor([-0.0979, -0.1447, 0.2425]),)
```



下一课时

链式法则

Thank You.
